



Search Keras documentation...

» [Keras API reference](#) / [Layers API](#) / [Attention layers](#) / MultiHeadAttention layer

MultiHeadAttention layer

MultiHeadAttention class

[source]

```
tf.keras.layers.MultiHeadAttention(  
    num_heads,  
    key_dim,  
    value_dim=None,  
    dropout=0.0,  
    use_bias=True,  
    output_shape=None,  
    attention_axes=None,  
    kernel_initializer="glorot_uniform",  
    bias_initializer="zeros",  
    kernel_regularizer=None,  
    bias_regularizer=None,  
    activity_regularizer=None,  
    kernel_constraint=None,  
    bias_constraint=None,  
    **kwargs  
)
```

MultiHeadAttention layer.

This is an implementation of multi-headed attention as described in the paper "Attention is all you Need" (Vaswani et al., 2017). If **query**, **key**, **value** are the same, then this is self-attention. Each timestep in **query** attends to the corresponding sequence in **key**, and returns a fixed-width vector.

This layer first projects **query**, **key** and **value**. These are (effectively) a list of tensors of length **num_attention_heads**, where the corresponding shapes are **(batch_size, <query dimensions>, key_dim)**, **(batch_size, <key/value dimensions>, key_dim)**, **(batch_size, <key/value dimensions>, value_dim)**.

Then, the query and key tensors are dot-producted and scaled. These are softmaxed to obtain attention probabilities. The value tensors are then interpolated by these probabilities, then concatenated back to a single tensor.

Finally, the result tensor with the last dimension as value_dim can take an linear projection and return.

When using MultiHeadAttention inside a custom Layer, the custom Layer must implement **build()** and call MultiHeadAttention's **_build_from_signature()**. This enables weights to be restored correctly when the model is loaded. TODO(b/172609172): link to documentation about calling custom build functions when used in a custom Layer.

Examples

Performs 1D cross-attention over two sequence inputs with an attention mask. Returns the additional attention weights over heads.

```
>>> layer = MultiHeadAttention(num_heads=2, key_dim=2)  
>>> target = tf.keras.Input(shape=[8, 16])  
>>> source = tf.keras.Input(shape=[4, 16])  
>>> output_tensor, weights = layer(target, source,  
...                               return_attention_scores=True)  
>>> print(output_tensor.shape)  
(None, 8, 16)  
>>> print(weights.shape)  
(None, 2, 8, 4)
```

Performs 2D self-attention over a 5D input tensor on axes 2 and 3.

```
>>> layer = MultiHeadAttention(num_heads=2, key_dim=2, attention_axes=(2, 3))
>>> input_tensor = tf.keras.Input(shape=[5, 3, 4, 16])
>>> output_tensor = layer(input_tensor, input_tensor)
>>> print(output_tensor.shape)
(None, 5, 3, 4, 16)
```

MultiHeadAttention layer

MultiHeadAttention class

Arguments

- **num_heads:** Number of attention heads.
- **key_dim:** Size of each attention head for query and key.
- **value_dim:** Size of each attention head for value.
- **dropout:** Dropout probability.
- **use_bias:** Boolean, whether the dense layers use bias vectors/matrices.
- **output_shape:** The expected shape of an output tensor, besides the batch and sequence dims. If not specified, projects back to the key feature dim.
- **attention_axes:** axes over which the attention is applied. **None** means attention over all axes, but batch, heads, and features.
- **kernel_initializer:** Initializer for dense layer kernels.
- **bias_initializer:** Initializer for dense layer biases.
- **kernel_regularizer:** Regularizer for dense layer kernels.
- **bias_regularizer:** Regularizer for dense layer biases.
- **activity_regularizer:** Regularizer for dense layer activity.
- **kernel_constraint:** Constraint for dense layer kernels.
- **bias_constraint:** Constraint for dense layer kernels.

Call arguments

- **query:** Query **Tensor** of shape **(B, T, dim)**.
- **value:** Value **Tensor** of shape **(B, S, dim)**.
- **key:** Optional key **Tensor** of shape **(B, S, dim)**. If not given, will use **value** for both **key** and **value**, which is the most common case.
- **attention_mask:** a boolean mask of shape **(B, T, S)**, that prevents attention to certain positions. The boolean mask specifies which query elements can attend to which key elements, 1 indicates attention and 0 indicates no attention. Broadcasting can happen for the missing batch dimensions and the head dimension.
- **return_attention_scores:** A boolean to indicate whether the output should be **(attention_output, attention_scores)** if **True**, or **attention_output** if **False**. Defaults to **False**.
- **training:** Python boolean indicating whether the layer should behave in training mode (adding dropout) or in inference mode (no dropout). Defaults to either using the training mode of the parent layer/model, or **False** (inference) if there is no parent layer.

Returns

- **attention_output:** The result of the computation, of shape **(B, T, E)**, where **T** is for target sequence shapes and **E** is the query input last dimension if **output_shape** is **None**. Otherwise, the multi-head outputs are project to the shape specified by **output_shape**.
- **attention_scores:** [Optional] multi-head attention coefficients over attention axes.